

House prediction analysis

Predicting house prices ties directly to real estate, economics, and business decisions.

The features listed are standard variables from this dataset:

- CRIM: Per capita crime rate by town.
- ZN: Proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS: Proportion of non-retail business acres per town.
- CHAS: Charles River dummy variable (1 if tract bounds river; 0 otherwise).
- NOX: Nitric oxides concentration (parts per 10 million).
- RM: Average number of rooms per dwelling.
- AGE: Proportion of owner-occupied units built prior to 1940.
- DIS: Weighted distances to five Boston employment centers.
- RAD: Index of accessibility to radial highways.
- TAX: Full-value property-tax rate per \$10,000.
- PTRATIO: Pupil-teacher ratio by town.
- B: $1000(Bk - 0.63)^2$ where Bk is the proportion of Black people by town.
- LSTAT: % lower status of the population.
- MEDV: Median value of owner-occupied homes in \$1000s. (This is the target variable)

Import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
import os
```

Loading the dataset

```
In [2]: # Define column names based on Boston Housing dataset
columns = [
    'CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE',
    'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV'
]

# Load data
df = pd.read_csv('/home/user/Documents/House prediction/house Prediction Data Set.csv', delim_whitespace=True, names=columns)
```

In [3]: # Display the first 10 rows of the dataframe
df.head(10)

Out[3]:

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B |
|---|---------|------|-------|------|-------|-------|-------|--------|-----|-------|---------|--------|
| 0 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296.0 | 15.3 | 396.90 |
| 1 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242.0 | 17.8 | 396.90 |
| 2 | 0.02729 | 0.0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242.0 | 17.8 | 392.83 |
| 3 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222.0 | 18.7 | 394.63 |
| 4 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222.0 | 18.7 | 396.90 |
| 5 | 0.02985 | 0.0 | 2.18 | 0 | 0.458 | 6.430 | 58.7 | 6.0622 | 3 | 222.0 | 18.7 | 394.12 |
| 6 | 0.08829 | 12.5 | 7.87 | 0 | 0.524 | 6.012 | 66.6 | 5.5605 | 5 | 311.0 | 15.2 | 395.60 |
| 7 | 0.14455 | 12.5 | 7.87 | 0 | 0.524 | 6.172 | 96.1 | 5.9505 | 5 | 311.0 | 15.2 | 396.90 |
| 8 | 0.21124 | 12.5 | 7.87 | 0 | 0.524 | 5.631 | 100.0 | 6.0821 | 5 | 311.0 | 15.2 | 386.63 |
| 9 | 0.17004 | 12.5 | 7.87 | 0 | 0.524 | 6.004 | 85.9 | 6.5921 | 5 | 311.0 | 15.2 | 386.71 |



In [4]: # Display the column names of the dataframe
df.columns

Out[4]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV'],
dtype='object')

In [5]: # Display the info of the dataframe
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column    Non-Null Count  Dtype  
 --- 
 0   CRIM      506 non-null   float64
 1   ZN        506 non-null   float64
 2   INDUS     506 non-null   float64
 3   CHAS      506 non-null   int64  
 4   NOX       506 non-null   float64
 5   RM         506 non-null   float64
 6   AGE        506 non-null   float64
 7   DIS        506 non-null   float64
 8   RAD        506 non-null   int64  
 9   TAX        506 non-null   float64
 10  PTRATIO   506 non-null   float64
 11  B          506 non-null   float64
 12  LSTAT     506 non-null   float64
 13  MEDV      506 non-null   float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

In [6]: # Display the statistical summary of the dataframe
df.describe()

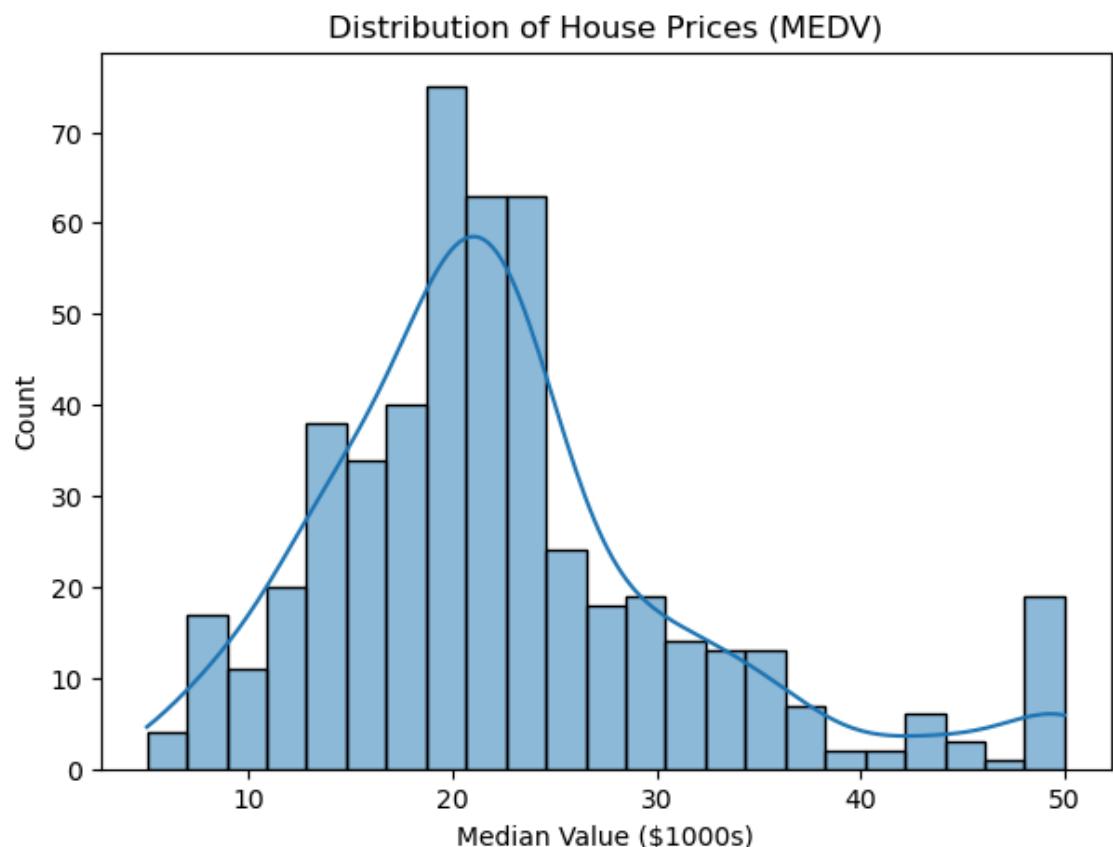
Out[6]:

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE |
|-------|------------|------------|------------|------------|------------|------------|------------|
| count | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 |
| mean | 3.613524 | 11.363636 | 11.136779 | 0.069170 | 0.554695 | 6.284634 | 68.574901 |
| std | 8.601545 | 23.322453 | 6.860353 | 0.253994 | 0.115878 | 0.702617 | 28.148861 |
| min | 0.006320 | 0.000000 | 0.460000 | 0.000000 | 0.385000 | 3.561000 | 2.900000 |
| 25% | 0.082045 | 0.000000 | 5.190000 | 0.000000 | 0.449000 | 5.885500 | 45.025000 |
| 50% | 0.256510 | 0.000000 | 9.690000 | 0.000000 | 0.538000 | 6.208500 | 77.500000 |
| 75% | 3.677083 | 12.500000 | 18.100000 | 0.000000 | 0.624000 | 6.623500 | 94.075000 |
| max | 88.976200 | 100.000000 | 27.740000 | 1.000000 | 0.871000 | 8.780000 | 100.000000 |

In [7]: # Drop duplicate rows from the dataframe
df.drop_duplicates(inplace=True)

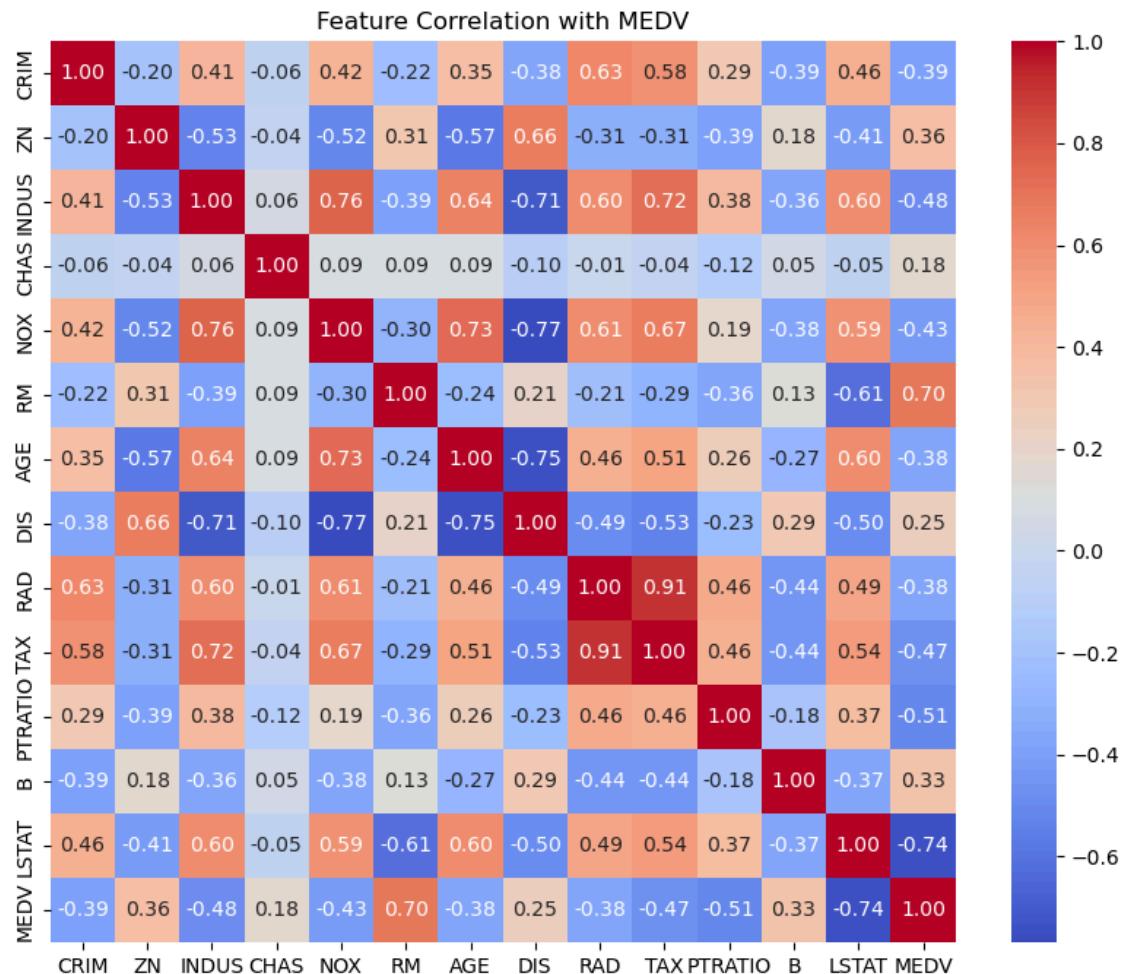
Data exploration

In [8]: # Plot the distribution of the target variable 'MEDV'
plt.figure(figsize=(7,5))
sns.histplot(df['MEDV'], kde=True)
plt.title('Distribution of House Prices (MEDV)')
plt.xlabel('Median Value (\$1000s)')
plt.show()



```
In [9]: # the data distribution of the house prices is right-skewed, indicating that
# there are more lower-priced houses than higher-priced ones.
```

```
In [10]: #correlation heatmap
plt.figure(figsize=(10,8))
corr = df.corr()
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Feature Correlation with MEDV')
plt.show()
```



```
In [11]: # Strongest Positive Correlation: RM (Average number of rooms) has a correlation of 0.70 with MEDV. This makes intuitive sense:
# houses with more rooms tend to be more valuable. Strongest Negative Correlation: LSTAT (% lower status of the population) has a correlation of -0.74 with MEDV. This also makes sense: areas with a higher percentage of lower-status residents tend to have lower home values. PTRATIO (Pupil-teacher ratio) is negatively correlated (-0.51), suggesting schools with fewer students per teacher are in more expensive areas. DIS (Distance to employment centers) is negatively correlated (-0.43), meaning homes closer to work are generally more expensive. RAD (Accessibility to highways) is positively correlated (0.46), indicating better highway access can increase value. NOX and INDUS are strongly positively correlated (0.76), which is logical as industrial areas often have higher pollution. TAX and PTRATIO are also positively correlated (0.91), suggesting towns with higher taxes might also have higher pupil-teacher ratios.
```

```
In [12]: X = df.drop('MEDV', axis=1)
y = df['MEDV']
```

```
In [13]: # Train-test split (80-20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Linear regression

```
In [14]: # Initialize and train
lr = LinearRegression()
lr.fit(X_train, y_train)
```

Out[14]:

▼ `LinearRegression` ⓘ ⓘ (https://scikit-learn.org/1.7/modules/generated/sklearn.linear_model.LinearRegression.html)

Parameters

```
In [15]: # Predict
y_pred_lr = lr.predict(X_test)
```

```
In [16]: # Evaluate
mse_lr = mean_squared_error(y_test, y_pred_lr)
rmse_lr = np.sqrt(mse_lr)
r2_lr = r2_score(y_test, y_pred_lr)
```

```
In [17]: print(f"Linear Regression Results:")
print(f"R²: {r2_lr:.4f}")
print(f"RMSE: {rmse_lr:.4f}")
```

Linear Regression Results:
R²: 0.6688
RMSE: 4.9286

```
In [18]: # The R2 value of 0.6688 indicates that approximately 66.88% of the
# variance in house prices can be explained by the features used
# in the model.
# RMSE of 4.9286 means that, on average, the model's predictions are
# off by about $4,928.60.
```

Random Forest classifier

```
In [19]: rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
```

```
In [20]: mse_rf = mean_squared_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mse_rf)
r2_rf = r2_score(y_test, y_pred_rf)
```

```
In [21]: print(f"\nRandom Forest Results:")
print(f"R²: {r2_rf:.4f}")
print(f"RMSE: {rmse_rf:.4f}")
```

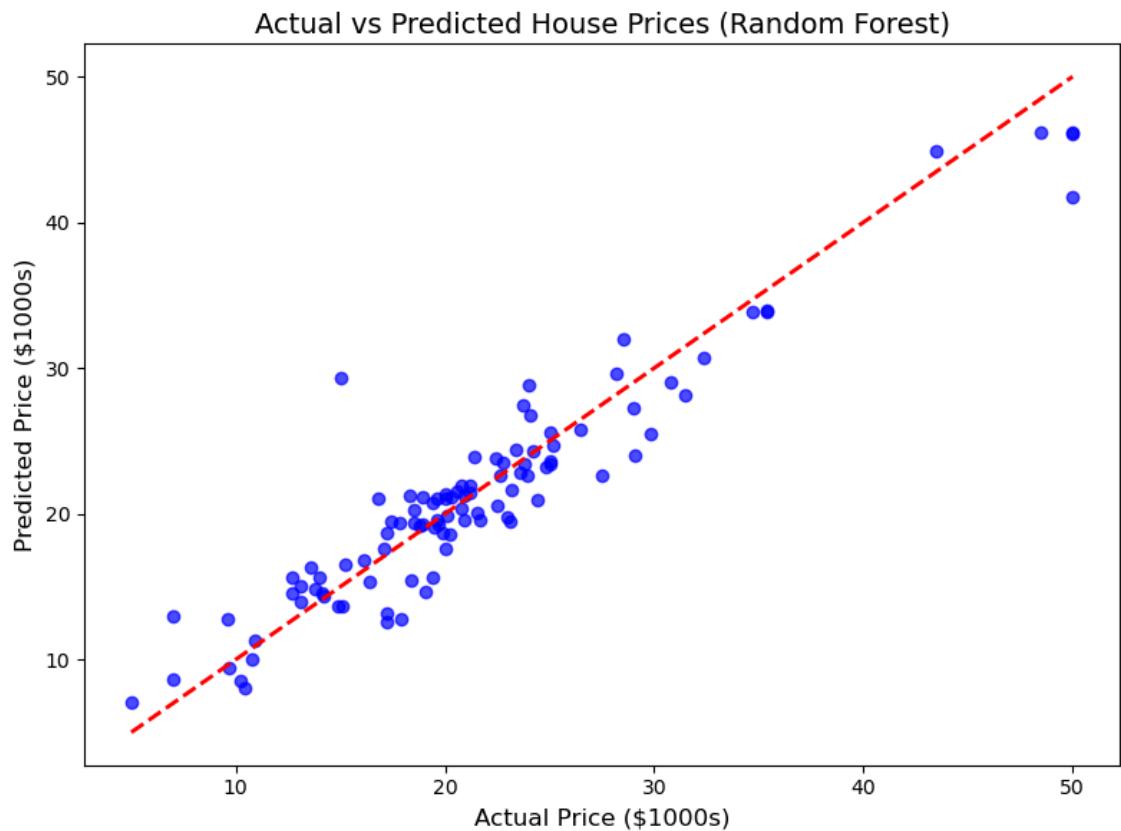
```
Random Forest Results:
R²: 0.8923
RMSE: 2.8110
```

```
In [22]: # The results from the random forest: R2 of 0.8923 indicate that about 89.23% of the variance in house prices is explained by the model.
# RMSE of 2.8110 indicate that on average the model's predictions are off by about $2,811.00. And the random forest model is outperforming
# the linear regression model significantly in both R2 and RMSE metrics.
```

Visualization

```
In [23]: os.makedirs('images', exist_ok=True)
```

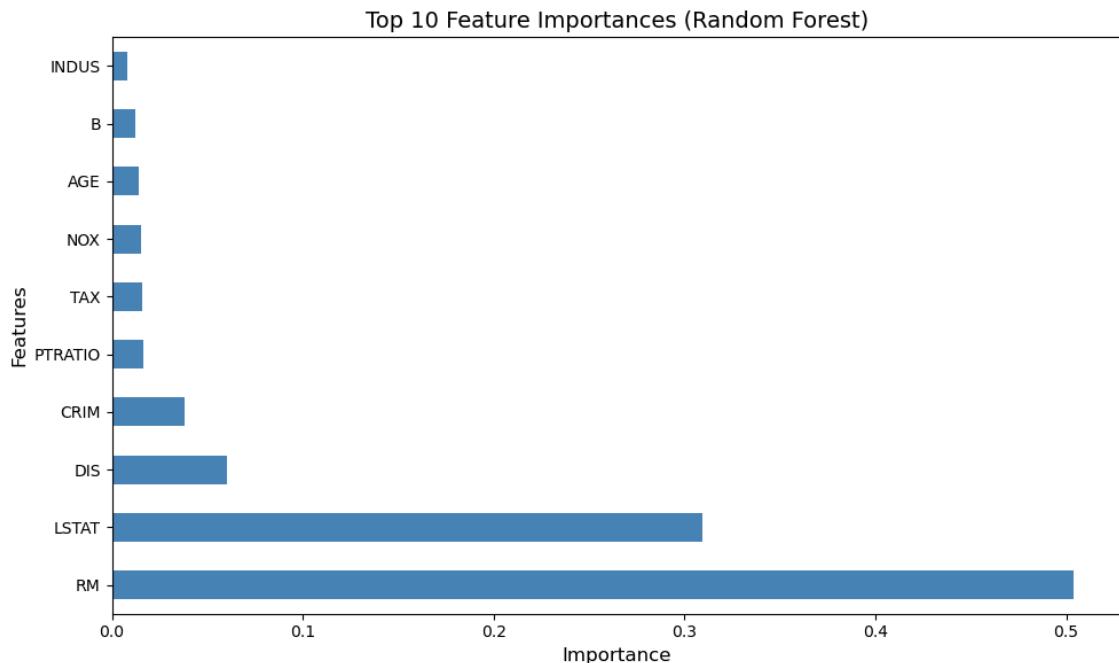
```
In [30]: # Plot Actual vs Predicted for Linear Regression
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred_rf, alpha=0.7, color='blue')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         'r--', lw=2)
plt.title('Actual vs Predicted House Prices (Random Forest)', fontsize=14)
plt.xlabel('Actual Price ($1000s)', fontsize=12)
plt.ylabel('Predicted Price ($1000s)', fontsize=12)
plt.tight_layout()
plt.savefig('images/actual_vs_predicted.png', dpi=300, bbox_inches='tight')
plt.show()
```



```
In [25]: # Most blue dots are clustered around the red dashed line – especially in the middle range (actual prices ~15–35).
# This means the model is doing a decent job predicting average-priced homes.
# For houses with actual prices > 40 (i.e., >$40,000), many predicted values fall below the red line.
# The model underestimates expensive homes.
# For very low actual prices (<10), some predictions are higher than they should be (dots above the red line).
# The model overestimates cheap homes.
# Overall, while the model captures general average trends, it struggles with extreme values on both ends.
```

```
In [31]: # Feature Importance from Random Forest
feat_importances = pd.Series(rf.feature_importances_, index=X.columns)
top_feats = feat_importances.nlargest(10)

plt.figure(figsize=(10, 6))
top_feats.plot(kind='barh', color='steelblue')
plt.title('Top 10 Feature Importances (Random Forest)', fontsize=14)
plt.xlabel('Importance', fontsize=12)
plt.ylabel('Features', fontsize=12)
plt.tight_layout()
plt.savefig('images/feature_importance_top10.png', dpi=300, bbox_inches='tight')
plt.show()
```



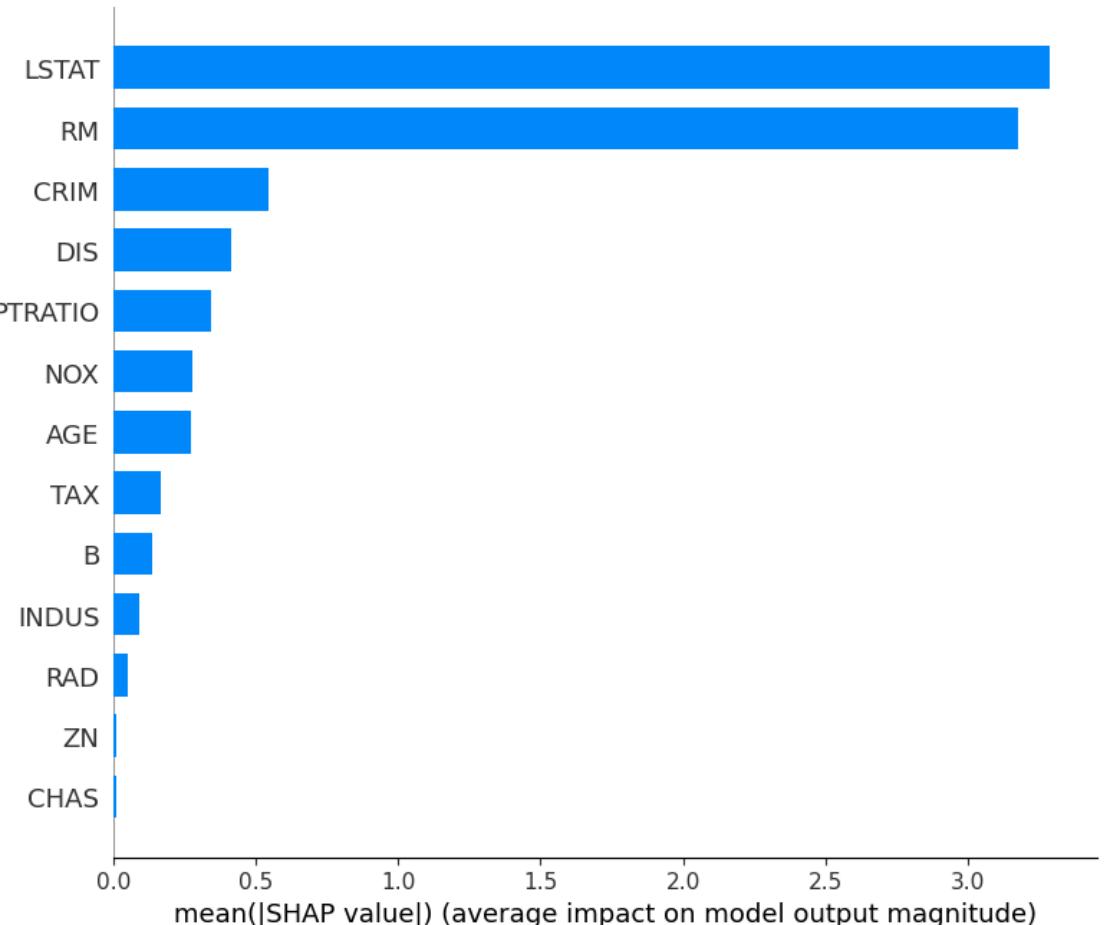
```
In [27]: # Of all the variables, the number of rooms per dwelling (RM) is the single biggest predictor of house price,
# followed by the percentage of lower-status residents (LSTAT). (DIS) Weighted distance to employment centers is the third most important feature. Other features play minor roles.
```

```
In [28]: #!pip install shap
```

```
In [29]: # Plot SHAP values for Random Forest
import shap

explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_test)

# Summary plot
shap.summary_plot(shap_values, X_test, plot_type="bar")
```



Key Insights from the Plot

Top 2 Most Influential Features

- LSTAT — Highest bar (~3.3+) % lower status of the population. Strong negative effect on house prices. The model uses this heavily to predict lower prices.
- RM — Second highest (~3.2) Average number of rooms per dwelling. Strong positive effect — more rooms → higher predicted price.

Moderately Important Features

- CRIM (crime rate) — ~0.6
- DIS (distance to employment centers) — ~0.5
- PTRATIO (pupil-teacher ratio) — ~0.4 These still have noticeable influence but less than LSTAT and RM.

Low-Impact Features

- NOX, AGE, TAX, B, INDUS — small bars (~0.2–0.3)
- RAD, ZN, CHAS — very tiny bars (<0.1) These contribute minimally to the model's predictions — possibly because they're redundant, noisy, or weakly correlated with MEDV.